

```

import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns

from sklearn.datasets import fetch_california_housing
housing = fetch_california_housing(as_frame = True)
print(housing)

{'data':      MedInc   HouseAge   AveRooms   AveBedrms   Population
AveOccup   Latitude \
0          8.3252    41.0     6.984127    1.023810    322.0    2.555556
37.88
1          8.3014    21.0     6.238137    0.971880    2401.0   2.109842
37.86
2          7.2574    52.0     8.288136    1.073446    496.0    2.802260
37.85
3          5.6431    52.0     5.817352    1.073059    558.0    2.547945
37.85
4          3.8462    52.0     6.281853    1.081081    565.0    2.181467
37.85
...
...
20635    1.5603    25.0     5.045455    1.133333    845.0    2.560606
39.48
20636    2.5568    18.0     6.114035    1.315789    356.0    3.122807
39.49
20637    1.7000    17.0     5.205543    1.120092    1007.0   2.325635
39.43
20638    1.8672    18.0     5.329513    1.171920    741.0    2.123209
39.43
20639    2.3886    16.0     5.254717    1.162264    1387.0   2.616981
39.37

      Longitude
0        -122.23
1        -122.22
2        -122.24
3        -122.25
4        -122.25
...
20635   -121.09
20636   -121.21
20637   -121.22
20638   -121.32
20639   -121.24

[20640 rows x 8 columns], 'target': 0           4.526
1           3.585
2           3.521

```

```

3      3.413
4      3.422
...
20635  0.781
20636  0.771
20637  0.923
20638  0.847
20639  0.894
Name: MedHouseVal, Length: 20640, dtype: float64, 'frame':
MedInc   HouseAge   AveRooms   AveBedrms   Population   AveOccup   Latitude
\
0       8.3252     41.0    6.984127    1.023810     322.0    2.555556
37.88
1       8.3014     21.0    6.238137    0.971880     2401.0   2.109842
37.86
2       7.2574     52.0    8.288136    1.073446     496.0    2.802260
37.85
3       5.6431     52.0    5.817352    1.073059     558.0    2.547945
37.85
4       3.8462     52.0    6.281853    1.081081     565.0    2.181467
37.85
...
...
20635  1.5603     25.0    5.045455    1.133333     845.0    2.560606
39.48
20636  2.5568     18.0    6.114035    1.315789     356.0    3.122807
39.49
20637  1.7000     17.0    5.205543    1.120092     1007.0   2.325635
39.43
20638  1.8672     18.0    5.329513    1.171920     741.0    2.123209
39.43
20639  2.3886     16.0    5.254717    1.162264     1387.0   2.616981
39.37

          Longitude  MedHouseVal
0           -122.23    4.526
1           -122.22    3.585
2           -122.24    3.521
3           -122.25    3.413
4           -122.25    3.422
...
20635      -121.09    0.781
20636      -121.21    0.771
20637      -121.22    0.923
20638      -121.32    0.847
20639      -121.24    0.894

```

[20640 rows x 9 columns], 'target_names': ['MedHouseVal'],
'feature_names': ['MedInc', 'HouseAge', 'AveRooms', 'AveBedrms',

```
'Population', 'AveOccup', 'Latitude', 'Longitude'], 'DESCR': '..\n_california_housing_dataset:\nCalifornia Housing dataset\n-----\n**Data Set Characteristics:**\n\n :Number of Instances: 20640\n :Number of Attributes: 8\n numeric, predictive attributes and the target\n\n :Attribute Information:\n - MedInc median income in block group\n - HouseAge median house age in block group\n - AveRooms average number of rooms per household\n - AveBedrms average number of bedrooms per household\n - Population block group population\n - AveOccup average number of household members\n - Latitude block group latitude\n - Longitude block group longitude\n\n :Missing Attribute Values:\nNone\n\nThis dataset was obtained from the StatLib repository.\n\nhttps://www.dcc.fc.up.pt/~ltorgo/Regression/cal\_housing.html\n\nThe target variable is the median house value for California districts,\nexpressed in hundreds of thousands of dollars ($100,000).\n\nThis dataset was derived from the 1990 U.S. census, using one row per\ncensus\nblock group. A block group is the smallest geographical unit\nfor which the U.S.\nCensus Bureau publishes sample data (a block group\ntypically has a population\nof 600 to 3,000 people).\n\nA household is\na group of people residing within a home. Since the average\nnumber of\nrooms and bedrooms in this dataset are provided per household, these\ncolumns may take surprisingly large values for block groups with few\nhouseholds\nand many empty houses, such as vacation resorts.\n\nIt can\nbe downloaded/loaded using the\n\n:func:`sklearn.datasets.fetch_california_housing` function.\n\n.. topic:: References\n - Pace, R. Kelley and Ronald Barry, Sparse\n Spatial Autoregressions,\n Statistics and Probability Letters, 33\n(1997) 291-297\n'}
```

```
housing['data'].head()
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
Latitude	8.3252 37.88	41.0	6.984127	1.023810	322.0	2.555556
Longitude	8.3014 37.86	21.0	6.238137	0.971880	2401.0	2.109842
MedInc	7.2574 37.85	52.0	8.288136	1.073446	496.0	2.802260
HouseAge	5.6431 37.85	52.0	5.817352	1.073059	558.0	2.547945
AveRooms	3.8462 37.85	52.0	6.281853	1.081081	565.0	2.181467
AveBedrms						
Population						
AveOccup						

```

3    -122.25
4    -122.25

housing['target'].head()

0    4.526
1    3.585
2    3.521
3    3.413
4    3.422
Name: MedHouseVal, dtype: float64

df = pd.DataFrame(housing['data'])
df

      MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup
Latitude \
0     8.3252      41.0   6.984127   1.023810      322.0  2.555556
37.88
1     8.3014      21.0   6.238137   0.971880     2401.0  2.109842
37.86
2     7.2574      52.0   8.288136   1.073446      496.0  2.802260
37.85
3     5.6431      52.0   5.817352   1.073059      558.0  2.547945
37.85
4     3.8462      52.0   6.281853   1.081081      565.0  2.181467
37.85
...
...
20635  1.5603      25.0   5.045455   1.133333      845.0  2.560606
39.48
20636  2.5568      18.0   6.114035   1.315789      356.0  3.122807
39.49
20637  1.7000      17.0   5.205543   1.120092     1007.0  2.325635
39.43
20638  1.8672      18.0   5.329513   1.171920      741.0  2.123209
39.43
20639  2.3886      16.0   5.254717   1.162264     1387.0  2.616981
39.37

      Longitude
0    -122.23
1    -122.22
2    -122.24
3    -122.25
4    -122.25
...
20635  -121.09
20636  -121.21
20637  -121.22

```

```
20638    -121.32
20639    -121.24
```

```
[20640 rows x 8 columns]
```

```
df['Price'] = housing['target']
df
```

```
      MedInc  HouseAge  AveRooms  AveBedrms  Population  AveOccup
Latitude \
0        8.3252     41.0   6.984127   1.023810      322.0  2.555556
37.88
1        8.3014     21.0   6.238137   0.971880     2401.0  2.109842
37.86
2        7.2574     52.0   8.288136   1.073446      496.0  2.802260
37.85
3        5.6431     52.0   5.817352   1.073059      558.0  2.547945
37.85
4        3.8462     52.0   6.281853   1.081081      565.0  2.181467
37.85
...
...
20635  1.5603     25.0   5.045455   1.133333      845.0  2.560606
39.48
20636  2.5568     18.0   6.114035   1.315789      356.0  3.122807
39.49
20637  1.7000     17.0   5.205543   1.120092     1007.0  2.325635
39.43
20638  1.8672     18.0   5.329513   1.171920      741.0  2.123209
39.43
20639  2.3886     16.0   5.254717   1.162264     1387.0  2.616981
39.37
```

```
      Longitude  Price
0        -122.23  4.526
1        -122.22  3.585
2        -122.24  3.521
3        -122.25  3.413
4        -122.25  3.422
...
20635    -121.09  0.781
20636    -121.21  0.771
20637    -121.22  0.923
20638    -121.32  0.847
20639    -121.24  0.894
```

```
[20640 rows x 9 columns]
```

```
df.info()
```

```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 20640 entries, 0 to 20639
Data columns (total 9 columns):
 #   Column      Non-Null Count  Dtype  
--- 
 0   MedInc       20640 non-null   float64
 1   HouseAge     20640 non-null   float64
 2   AveRooms     20640 non-null   float64
 3   AveBedrms    20640 non-null   float64
 4   Population    20640 non-null   float64
 5   AveOccup     20640 non-null   float64
 6   Latitude      20640 non-null   float64
 7   Longitude     20640 non-null   float64
 8   Price         20640 non-null   float64
dtypes: float64(9)
memory usage: 1.4 MB

```

```
df.isna().sum()
```

```

MedInc      0
HouseAge    0
AveRooms    0
AveBedrms   0
Population  0
AveOccup    0
Latitude    0
Longitude   0
Price       0
dtype: int64

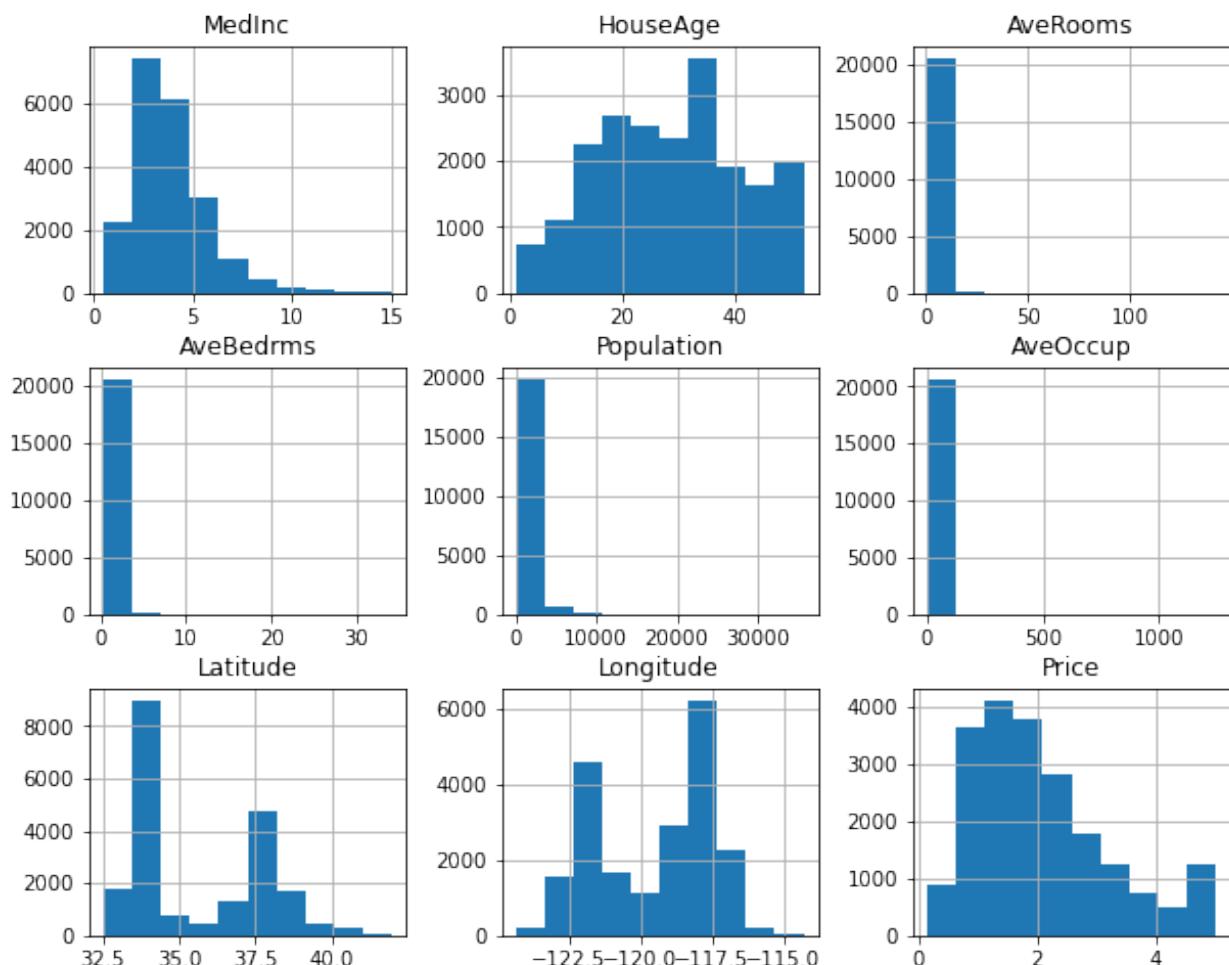
```

```
df.describe().T
```

	count	mean	std	min	25%
MedInc	20640.0	3.870671	1.899822	0.499900	2.563400
HouseAge	20640.0	28.639486	12.585558	1.000000	18.000000
AveRooms	20640.0	5.429000	2.474173	0.846154	4.440716
AveBedrms	20640.0	1.096675	0.473911	0.333333	1.006079
Population	20640.0	1425.476744	1132.462122	3.000000	787.000000
AveOccup	20640.0	3.070655	10.386050	0.692308	2.429741
Latitude	20640.0	35.631861	2.135952	32.540000	33.930000
Longitude	20640.0	-119.569704	2.003532	-124.350000	-121.800000
Price	20640.0	2.068558	1.153956	0.149990	1.196000

	50%	75%	max
MedInc	3.534800	4.743250	15.000100
HouseAge	29.000000	37.000000	52.000000
AveRooms	5.229129	6.052381	141.909091
AveBedrms	1.048780	1.099526	34.066667
Population	1166.000000	1725.000000	35682.000000
AveOccup	2.818116	3.282261	1243.333333
Latitude	34.260000	37.710000	41.950000
Longitude	-118.490000	-118.010000	-114.310000
Price	1.797000	2.647250	5.000010

```
df.hist(figsize=(10,8))
plt.show()
```



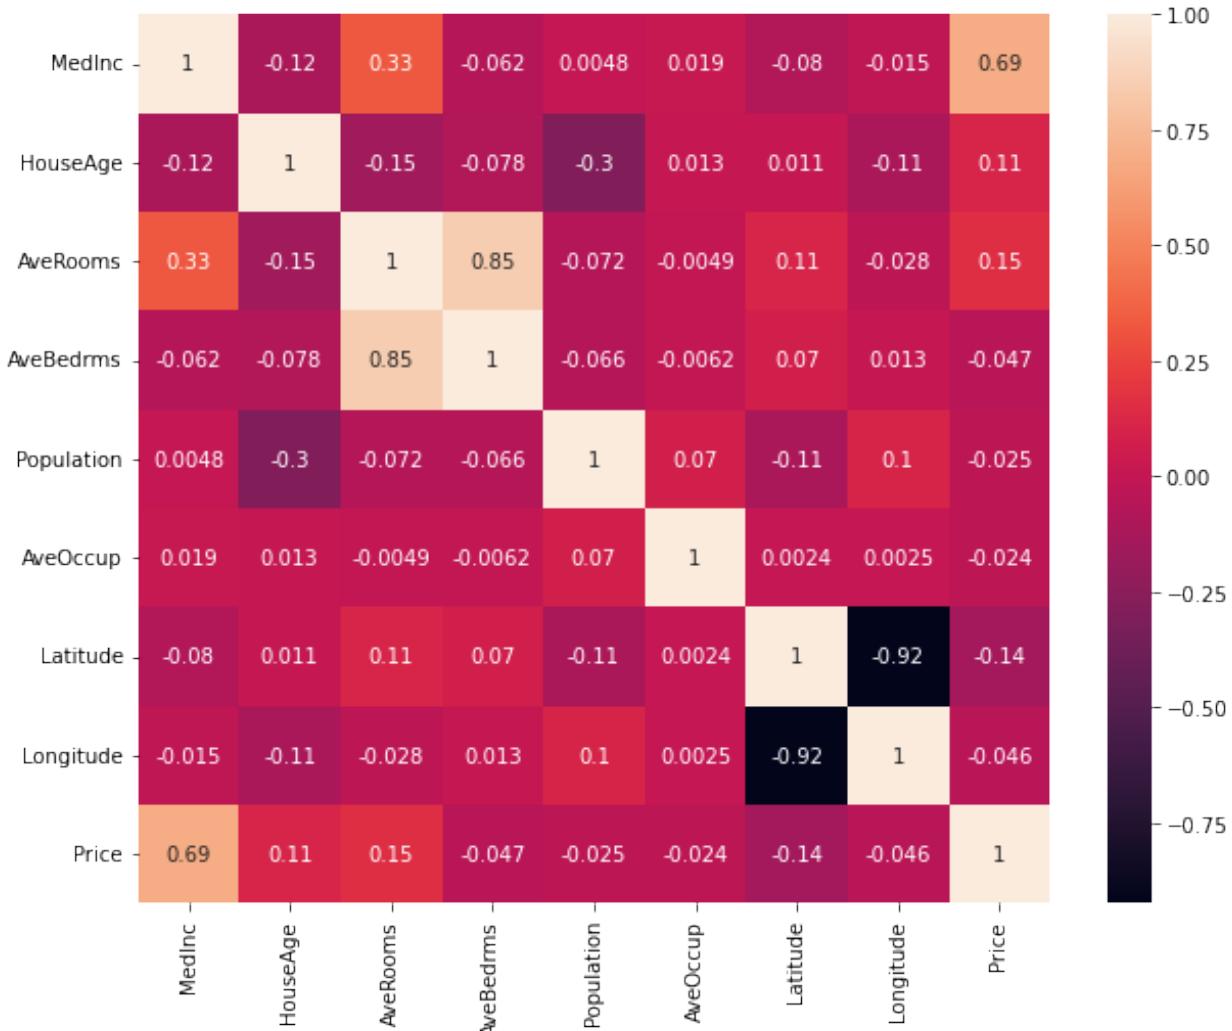
```
corr = df.corr()
corr
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population
AveOccup	\				
MedInc	1.000000	-0.119034	0.326895	-0.062040	0.004834
0.018766					
HouseAge	-0.119034	1.000000	-0.153277	-0.077747	-0.296244
0.013191					
AveRooms	0.326895	-0.153277	1.000000	0.847621	-0.072213
0.004852					
AveBedrms	-0.062040	-0.077747	0.847621	1.000000	-0.066197
0.006181					
Population	0.004834	-0.296244	-0.072213	-0.066197	1.000000
0.069863					
AveOccup	0.018766	0.013191	-0.004852	-0.006181	0.069863
1.000000					
Latitude	-0.079809	0.011173	0.106389	0.069721	-0.108785
0.002366					
Longitude	-0.015176	-0.108197	-0.027540	0.013344	0.099773
0.002476					
Price	0.688075	0.105623	0.151948	-0.046701	-0.024650
0.023737					

	Latitude	Longitude	Price
MedInc	-0.079809	-0.015176	0.688075
HouseAge	0.011173	-0.108197	0.105623
AveRooms	0.106389	-0.027540	0.151948
AveBedrms	0.069721	0.013344	-0.046701
Population	-0.108785	0.099773	-0.024650
AveOccup	0.002366	0.002476	-0.023737
Latitude	1.000000	-0.924664	-0.144160
Longitude	-0.924664	1.000000	-0.045967
Price	-0.144160	-0.045967	1.000000

```
plt.figure(figsize=(10,8))
sns.heatmap(corr, annot=True)
```

<AxesSubplot:>



```
X = df.drop('Price',axis=1)  
y = df['Price']
```

X

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
Latitude \ 0	8.3252	41.0	6.984127	1.023810	322.0	2.555556
37.88						
1	8.3014	21.0	6.238137	0.971880	2401.0	2.109842
37.86						
2	7.2574	52.0	8.288136	1.073446	496.0	2.802260
37.85						
3	5.6431	52.0	5.817352	1.073059	558.0	2.547945
37.85						
4	3.8462	52.0	6.281853	1.081081	565.0	2.181467
37.85						
...

```
...
20635    1.5603      25.0   5.045455   1.133333     845.0   2.560606
39.48
20636    2.5568      18.0   6.114035   1.315789     356.0   3.122807
39.49
20637    1.7000      17.0   5.205543   1.120092    1007.0   2.325635
39.43
20638    1.8672      18.0   5.329513   1.171920     741.0   2.123209
39.43
20639    2.3886      16.0   5.254717   1.162264    1387.0   2.616981
39.37
```

```
        Longitude
0          -122.23
1          -122.22
2          -122.24
3          -122.25
4          -122.25
...
20635      ...
20636      -121.09
20637      -121.21
20638      -121.22
20639      -121.32
```

```
[20640 rows x 8 columns]
```

```
y
```

```
0      4.526
1      3.585
2      3.521
3      3.413
4      3.422
```

```
...
20635      0.781
20636      0.771
20637      0.923
20638      0.847
20639      0.894
```

```
Name: Price, Length: 20640, dtype: float64
```

```
from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X,y, test_size =
0.25,random_state = 42)
```

```
df.shape
```

```
(20640, 9)
```

```
X_train
```

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
Latitude \						
8158	4.2143	37.0	5.288235	0.973529	860.0	2.529412
33.81						
18368	5.3468	42.0	6.364322	1.087940	957.0	2.404523
37.16						
19197	3.9191	36.0	6.110063	1.059748	711.0	2.235849
38.45						
3746	6.3703	32.0	6.000000	0.990196	1159.0	2.272549
34.16						
13073	2.3684	17.0	4.795858	1.035503	706.0	2.088757
38.57						
...
...						
11284	6.3700	35.0	6.129032	0.926267	658.0	3.032258
33.78						
11964	3.0500	33.0	6.868597	1.269488	1753.0	3.904232
34.02						
5390	2.9344	36.0	3.986717	1.079696	1756.0	3.332068
34.03						
860	5.7192	15.0	6.395349	1.067979	1777.0	3.178891
37.58						
15795	2.5755	52.0	3.402576	1.058776	2619.0	2.108696
37.77						
Longitude						
8158	-118.12					
18368	-121.98					
19197	-122.69					
3746	-118.41					
13073	-121.33					
...	...					
11284	-117.96					
11964	-117.43					
5390	-118.38					
860	-121.96					
15795	-122.42					

[15480 rows x 8 columns]

X_test

	MedInc	HouseAge	AveRooms	AveBedrms	Population	AveOccup
Latitude \						
20046	1.6812	25.0	4.192201	1.022284	1392.0	3.877437
36.06						
3024	2.5313	30.0	5.039384	1.193493	1565.0	2.679795
35.14						
15663	3.4801	52.0	3.977155	1.185877	1310.0	1.360332
37.80						

```
20484 5.7376      17.0 6.163636   1.020202      1705.0 3.444444
34.28
9814  3.7250      34.0 5.492991   1.028037      1063.0 2.483645
36.62
...
...
5363  6.6260      51.0 5.532213   0.974790      771.0 2.159664
34.04
19755 2.1898      30.0 4.509091   0.945455      410.0 2.484848
40.18
4885  2.1667      37.0 3.272152   1.056962      2173.0 4.584388
34.02
13043 6.8869      6.0  7.382385   1.030075      2354.0 2.528464
38.51
8583  6.6321      36.0 5.734644   1.056511      1033.0 2.538084
33.89
```

```
    Longitude
20046 -119.01
3024  -119.46
15663 -122.44
20484 -118.72
9814  -121.93
...
...
5363  -118.42
19755 -122.21
4885  -118.26
13043 -121.06
8583  -118.40
```

```
[5160 rows x 8 columns]
```

```
y_train.shape
```

```
(15480,)
```

```
y_test.shape
```

```
(5160,)
```

```
from sklearn.linear_model import LinearRegression
model = LinearRegression()
model.fit(X_train,y_train)
```

```
LinearRegression()
```

```
y_pred = model.predict(X_test)
y_pred
```

```

array([0.72412832, 1.76677807, 2.71151581, ..., 1.72382152,
2.34689276,
      3.52917352])

from sklearn.metrics import mean_squared_error, mean_absolute_error,
r2_score
print(mean_squared_error(y_pred,y_test))

0.541128747847068

print(mean_absolute_error(y_pred,y_test))

0.529696401291945

r2_score(y_test,y_pred)

0.5910509795491358

```

Make Prediction

```

new_data = [[7.2574, 52.0, 8.288136, 1.073446, 496.0, 2.802260, 37.85, -
122.24]]
model.predict(new_data)

C:\Users\piyus\anaconda3\lib\site-packages\sklearn\base.py:464:
UserWarning: X does not have valid feature names, but LinearRegression
was fitted with feature names
    warnings.warn(
array([3.66368929])

```

MSE

1. Mean Squared Error (MSE): MSE is a measure of the average squared difference between the predicted values and the actual values. It gives more weight to larger errors. The formula for MSE is:

$$\text{MSE} = (1/n) * \sum (y_i - \hat{y}_i)^2$$

Where:

- n is the number of data points.
- y_i represents the actual values.
- \hat{y}_i represents the predicted values.

Let's illustrate this with an example:

Actual values (y_i): [2, 3, 5, 7, 10] Predicted values (\hat{y}_i): [1.5, 2.5, 4.5, 7.5, 9.5]

$$\text{MSE} = (1/5) * [(2 - 1.5)^2 + (3 - 2.5)^2 + (5 - 4.5)^2 + (7 - 7.5)^2 + (10 - 9.5)^2]$$

$$\text{MSE} = (1/5) * [0.25 + 0.25 + 0.25 + 0.25 + 0.25] \text{ MSE} = 0.25$$

```

import numpy as np
y_actual = np.array([2, 3, 5, 7, 10])
y_pred = np.array([1.5, 2.5, 4.5, 7.5, 9.5])
print(mean_squared_error(y_actual,y_pred))

```

0.25

MAE

- Mean Absolute Error (MAE): MAE is a measure of the average absolute difference between the predicted values and the actual values. It gives equal weight to all errors. The formula for MAE is:

$$\text{MAE} = (1/n) * \sum |y_i - \hat{y}_i|$$

Using the same example:

$$\text{MAE} = (1/5) * [|2 - 1.5| + |3 - 2.5| + |5 - 4.5| + |7 - 7.5| + |10 - 9.5|] \text{MAE} = (1/5) * [0.5 + 0.5 + 0.5 + 0.5] \text{MAE} = 0.5$$

```

print(mean_absolute_error(y_actual,y_pred))

```

0.5

R-squared (R^2) Score

- R-squared (R^2) Score: R-squared is a measure of how well the regression model fits the data. It represents the proportion of the variance in the dependent variable (y) that is explained by the independent variables (X). The formula for R^2 is:

$$R^2 = 1 - (\text{MSE(model)} / \text{MSE(mean)})$$

Where:

- MSE(model) is the mean squared error of the model.
- MSE(mean) is the mean squared error of the mean of the actual values.

In this example, we already calculated MSE as 0.25. Let's assume the mean of the actual values is 5:

$$\begin{aligned} \text{MSE(mean)} &= (1/5) * [(2 - 5)^2 + (3 - 5)^2 + (5 - 5)^2 + (7 - 5)^2 + (10 - 5)^2] \\ \text{MSE(mean)} &= (1/5) * [9 + 4 + 0 + 4 + 25] \text{MSE(mean)} = 8.4 \end{aligned}$$

Now, calculate R^2 :

$R^2 = 1 - (0.25 / 8.4)$ $R^2 \approx 0.9702$ An R^2 score close to 1 indicates that the model explains a high proportion of the variance in the data, while lower values suggest

that the model does not fit the data well. In this example, the R² score of approximately 0.9702 indicates a good fit between the model and the data.

```
r2_score(y_actual,y_pred)  
0.9696601941747572
```